

# *DESIGN A KAFKA APPLICATION*

Event Streaming with Kafka



## Table of Contents

<b><i>Exercise 1 - Design a Kafka-based IoT Data Pipeline.....</i></b>	<b><i>3</i></b>
Step 1: Identify IoT devices and data streams .....	3
Step 2: Define Kafka topics .....	3
Step 3: Design the data ingestion layer .....	3
Step 4: Design the Kafka cluster.....	3
Step 5: Design the data processing and analysis layer .....	3
Step 6: Design the data storage layer.....	3
Step 7: Design the visualization and alerting layer .....	3
Step 8: Sketch the architecture .....	4
<b><i>Exercise 2 - Design a Kafka-based Log Aggregation System .....</i></b>	<b><i>5</i></b>
Step 1: Identify log sources.....	5
Step 2: Define Kafka topics .....	5
Step 3: Design the log ingestion layer .....	5
Step 4: Design the Kafka cluster.....	5
Step 5: Design the log processing and analysis layer .....	5
Step 6: Design the log storage layer .....	5
Step 7: Design the monitoring and alerting layer .....	6
Step 8: Sketch the architecture .....	6
<b><i>Challenge: Exercise 3 – Design a Kafka Microservices Architecture .....</i></b>	<b><i>7</i></b>
Step 1: Identify microservices.....	7
Step 2: Define Kafka topics .....	7
Step 3: Design the Kafka cluster.....	7
Step 4: Design Kafka producers and consumers for microservices.....	7
Step 5: Design external data storage .....	7
Step 6: Design load balancers and API gateways.....	7
Step 7: Sketch the architecture .....	8

## Exercise 1 - Design a Kafka-based IoT Data Pipeline

**Objective:** Design an architecture for a smart city with multiple IoT devices using Kafka to process and analyze the data generated by these devices in real-time.

**Scenario:** A smart city with various IoT devices, such as traffic sensors, weather stations, and smart energy meters, generates a massive volume of data. You must create an architecture that processes, analyzes, and stores this data efficiently using Kafka.

### Step 1: Identify IoT devices and data streams

List the IoT devices in the smart city scenario, such as traffic sensors, weather stations, and smart energy meters. Determine the data streams generated by each device.

### Step 2: Define Kafka topics

Create Kafka topics for each data stream or group of related data streams. Consider organizing topics by data types, devices, or functional areas.

### Step 3: Design the data ingestion layer

Plan how IoT devices will send data to Kafka as producers. Determine the connection and communication protocols to be used.

### Step 4: Design the Kafka cluster

Determine the number of Kafka brokers needed for the data volume and throughput requirements. Plan for replication and partitioning strategies for fault-tolerance and scalability.

### Step 5: Design the data processing and analysis layer

Identify the types of processing and analysis needed for each data stream, such as data aggregation, anomaly detection, or insight generation. Design Kafka consumers to perform these tasks and consume data from the appropriate topics.

### Step 6: Design the data storage layer

Choose the appropriate data storage systems for processed data, such as databases or data lakes. Plan how Kafka consumers will write processed data to these storage systems.

### Step 7: Design the visualization and alerting layer

Plan a real-time dashboard for visualizing data and insights. Design an alerting and notification system for critical events.

## Step 8: Sketch the architecture

Draw the architecture, showing the flow of data from IoT devices to the final storage and visualization components. Clearly label each component and the Kafka topics involved in the data pipeline. You can do this with any tool you like, for example, draw.io.

## Exercise 2 - Design a Kafka-based Log Aggregation System

Objective: Design an architecture for a distributed web application using Kafka to aggregate, process, and analyze logs generated by multiple services.

Scenario: A distributed web application consists of multiple services, such as web servers, application servers, and databases, each generating logs. You must create an architecture that collects logs from all services, processes them using Kafka, and feeds the processed data into a monitoring and alerting system. The goal is to efficiently aggregate and analyze logs for better application performance and troubleshooting.

### Step 1: Identify log sources

List the services in the distributed web application, such as web servers, application servers, and databases. Determine the types of logs generated by each service (e.g., access logs, error logs, transaction logs).

### Step 2: Define Kafka topics

Create Kafka topics for different log types or groups of related log types. Consider organizing topics by log types, services, or functional areas.

### Step 3: Design the log ingestion layer

Plan how logs from each service will be sent to Kafka as producers. Determine the connection and communication protocols to be used, such as log shippers or agents.

### Step 4: Design the Kafka cluster

Determine the number of Kafka brokers needed for the log volume and throughput requirements. Plan for replication and partitioning strategies for fault-tolerance and scalability.

### Step 5: Design the log processing and analysis layer

Identify the types of processing and analysis needed for each log type, such as log parsing, data enrichment, anomaly detection, or pattern recognition. Design Kafka consumers to perform these tasks and consume data from the appropriate topics.

### Step 6: Design the log storage layer

Choose the appropriate storage systems for processed log data, such as databases or data lakes. Plan how Kafka consumers will write processed log data to these storage systems.

### Step 7: Design the monitoring and alerting layer

Plan a monitoring system for visualizing log data, detecting anomalies, and providing insights. Design an alerting and notification system for critical events or incidents.

### Step 8: Sketch the architecture

Draw the architecture, showing the flow of logs from the services to the final storage and monitoring components. Clearly label each component and the Kafka topics involved in the log aggregation pipeline.

## Challenge: Exercise 3 – Design a Kafka Microservices Architecture

Objective: Design a microservices-based e-commerce system using Kafka for communication between services.

Scenario: An e-commerce system with multiple microservices, such as user management, product catalog, shopping cart, and order processing. You should create an architecture illustrating how these microservices interact and communicate with each other through Kafka topics.

### Step 1: Identify microservices

List the microservices in the e-commerce system, such as user management, product catalog, shopping cart, and order processing.

### Step 2: Define Kafka topics

Create Kafka topics for different communication streams between microservices, such as user events, product updates, cart events, and order events. Consider organizing topics based on functional areas, data types, or events.

### Step 3: Design the Kafka cluster

Determine the number of Kafka brokers needed for the data volume and throughput requirements. Plan for replication and partitioning strategies for fault-tolerance and scalability.

### Step 4: Design Kafka producers and consumers for microservices

Plan how each microservice will send and receive messages using Kafka producers and consumers. Determine how each microservice will interact with Kafka topics for sending and receiving messages.

### Step 5: Design external data storage

Choose the appropriate data storage systems for each microservice, such as databases. Plan how microservices will interact with these storage systems to store service-specific data.

### Step 6: Design load balancers and API gateways

Plan how external requests to the system will be handled using load balancers and API gateways. Determine the routing and security rules for handling requests to various microservices.



### Step 7: Sketch the architecture

Draw the architecture, showing the microservices, Kafka components, and data flow between them. Clearly label each microservice and the Kafka topics involved in the communication.